

# Six Levels of Privacy: A Framework for Financial Synthetic Data

Tucker Balch, Vamsi K. Potluru, Deepak Paramanand, Manuela Veloso

J.P. Morgan Chase & Company AI Research  
January 29, 2024

## Abstract

Synthetic Data is increasingly important in financial applications. In addition to the benefits it provides, such as improved financial modeling and better testing procedures, it poses privacy risks as well. Such data may arise from client information, business information, or other proprietary sources that must be protected. Even though the process by which Synthetic Data is generated serves to obscure the original data to some degree, the extent to which privacy is preserved is hard to assess. Accordingly, we introduce a hierarchy of “levels” of privacy that are useful for categorizing Synthetic Data generation methods and the progressively improved protections they offer. While the six levels were devised in the context of financial applications, they may also be appropriate for other industries as well. Our paper includes: A brief overview of Financial Synthetic Data, how it can be used, how its value can be assessed, privacy risks, and privacy attacks. We close with details of the “Six Levels” that include defenses against those attacks.

## 1 Introduction

As the name suggests, Synthetic Data is artificially generated rather than produced by real world events. Synthetic Data is created via two primary methods, namely: 1) By *transforming* real data, or 2) By *simulation* of real processes. We refer the reader to the rich literature on Synthetic Data and the many mechanisms for creating it [ADM<sup>+</sup>20]. In financial applications we focus on three key uses for Synthetic Data:

1. **Liberate data:** Depending on its source, the sensitivity or risk associated with particular types of data can be significantly reduced or eliminated when transformed to synthetic form. We might be able to, accordingly, share it more freely and with less risk. We refer to this as “liberating data.”

2. **Augment for training:** Synthetic Data can be used to augment real data used for training models in order to fill gaps in the coverage of the data. In some cases the models trained in this way perform better than those without augmentation.
3. **Testing:** With Synthetic Data we have the advantage of being able to control the generation so that we know its properties and contents. If for instance, we want to test a fraud detection algorithm, we can “plant” known fraudulent patterns in the data to check if an algorithm flags them. Synthetic Data can also be used to explore the “corner cases” to see of the processes that use the data break under stress.

The value of Synthetic Data for each of the above uses may vary according to the application. Three different properties of the data contribute to an assessment of its value. As you will see, these properties are sometimes confounding: It is usually not possible for a dataset to score well along all three dimensions at once. The dimensions include:

- **Realism:** How realistic is the data, in the sense that it matches the real process or dataset that we seek to emulate? In general, the higher the fidelity of the Synthetic Data, the more useful it is for downstream processes, but at the cost of reduced privacy.
- **Privacy:** How easy is it for an adversary to “reverse engineer” the dataset to infer properties of the original data? In some cases, it is possible to discover specific private information about individual records in the original data even though they are not present in the Synthetic Data (see Section 2: Privacy Attacks). Other proprietary or competitive information might also be revealed such as the distributional properties of data elements like age, salary, or credit rating of a client list.
- **Utility:** How well does the data serve the purpose for which it was created? As one example, we might want to use the data to augment real data in the training of an ML model. We would evaluate utility in this case by measuring the uplift the data provides for the model: E.g., Are its predictions now more accurate? In another case, we might be using the data to test an existing model or process, say for processing credit card transactions. These tests might be aimed at discovering “breaks” in the data processing pipeline (e.g., are large, or negative transactions handled appropriately?)

The metrics are interrelated, for instance: Increased realism usually suggests reduced privacy; Increased privacy may degrade utility. Note that while one might assume realism is the most important factor, this is not always the case. If we are testing a product or process and we only use real, or historical data, we might not expose flaws regarding how the system would respond to new, unexpected scenarios.

In the next section we consider some of the risks regarding privacy for financial data.

## 2 Privacy risks for financial data

Financial institutions are appropriately protective of their data and the data they hold for their clients. Data sharing between various lines of business within a company, and potentially, externally with clients or vendors, is governed by regulations and internal guidelines. These controls are designed to protect clients' sensitive information and protect firms from the unauthorized sharing of MNPI (Material Non-Public Information), as well as litigation, reputation, and competitive risks.

Here we review some prominent risks and relevant regulations that apply to financial institutions. While specific to this industry, these regulations are representative of those many businesses face.

- **Fair Credit Reporting Act (FCRA):**

This U.S. law requires that information collected by consumer reporting agencies (e.g. credit bureaus) cannot be provided to anyone who does not have a purpose specified in the Act. In particular, the data cannot be used for other purposes even if data that identify an individual are removed. In addition, the data user must ensure that identity cannot be inferred using other non-Personally Identifiable Information (PII) data fields.

- **Regulation on Unfair, Deceptive or Abusive Acts or Practices (UDAAP):** In many cases consumers and clients can specify how their personal data can be used. Sharing such data is a UDAAP violation if used or shared in a manner contrary to the choices made by, or representations made to, consumers or clients. In particular, in many settings data is subject to privacy elections made by consumers.

- **Litigation risks:** Inappropriate release of data or functions of data (e.g., models trained on data, insights from data, or synthetic data resembling these datasets) that reveal PII or statistics (e.g., global characteristics) of the data, may pose litigation risks. This is particularly important in the context of data sourced from external vendors: Use of such data is typically constrained by contracts that precisely define the scope of the use.

- **Competitive risks:** Publishing data that reveals the characteristics of a firm's client base or industries and publicly traded companies the firm has interest in, may pose competitive, antitrust and increased insider trading risks. This might apply even if the published data is synthetic.

## 3 Privacy attacks

In order to appropriately assess the protections privacy measures might provide, we must consider how data might be exploited or "attacked" by an adversary [SZZ<sup>+</sup>23]. We assume there exists an adversary who aims to extract private

	FCRA	UDAAP	Litigation Risk	Competitive Risk
Membership Inference Attack	Applicable	Applicable	Applicable	N/A
Attribute Inference Attack	Applicable	Applicable	Applicable	N/A
Property Inference Attack	N/A	N/A	Applicable	Applicable
Model Inference Attack	Applicable	Applicable	Applicable	Applicable

Table 1: Privacy attacks on synthetic data can lead to breach of various regulations in financial applications.

information from Synthetic Data or from some other output model output. Each type of attack is characterized by assumptions including: What information is available to the adversary? What information should be protected? What is the goal of the attack? Here we enumerate the most relevant attacks. Also, see Table 1 for an analysis of attacks versus regulatory risks.

- **Reconstruction attacks** Also known as attribute inference attacks. Reconstruction attacks are characterized by an adversary in possession of partial knowledge of a set of features with the aim to recover *sensitive* features or the full data sample. For example, if some columns matching public information for an individual (e.g. from voter registration data) correspond to an entry in the candidate dataset that also includes private attributes (e.g. credit card billing records), the presence of the individual can reveal the values of the private attributes for that person. [NS07].
- **Membership inference attacks (MIAs)** In many cases the presence of an individual’s data in a dataset by itself can reveal sensitive information. The adversary’s task in MIA is to infer whether an individual was present in the training dataset or not [SSS16]. An adversary with knowledge of an individual’s presence in the dataset can further exploit that knowledge in linkage (or reconstruction) attacks to identify sensitive attributes of that individual. Thus, MIA can be used as a stepping stone to launch other types of attack.
- **Property inference attacks** Property inference represents the ability to extract properties of the original dataset from the corresponding synthetic data. In general, property inference refers to learning summary statistics of the original data (e.g. mean value, quantiles, histograms etc.) under the assumption of access to Synthetic Data only. Note that preventing property inference attacks necessarily degrades fidelity of the synthetic data [LWSF23].

## 4 Privacy levels

Now we introduce a six-level privacy defense hierarchy and discuss the privacy attacks, utility implications, and potential privacy guarantees for each level. Each level corresponds to a group of defense mechanisms with increasingly stronger privacy protections.

These levels can provide guidance to businesses regarding the security and utility of their Synthetic Data. For instance, they might choose to allow internal sharing of Level 2 data if it arises from a non-critical source, but require Level 4 protections for more sensitive data. The relevant privacy level should be determined according to the use case to balance multiple objectives such as the business goal, security, speed of generation, and utility.

In the first 4 levels, we consider methods where the data is *transformed* from the original dataset to the Synthetic Data. In the figures, the original data appears on the left, and the arrows indicate how the data is transformed. We focus on tabular data in these examples, but the principles can apply to other types of data.




Figure 1: Privacy Level 1: Obscure PII

### 4.1 Privacy Level 1: Obscure PII

Examples of mechanisms at this level include dropping, replacing, masking, or anonymizing the PII attributes. Since this approach does not modify non-PII attributes in any way, it does not reduce the utility of downstream tasks and accordingly there is no utility degradation. This however represents weak privacy protection as data remains vulnerable to reconstruction attacks [NS07].




Figure 2: Privacy Level 2: Obscure PII + noise

## 4.2 Privacy Level 2: Obscure PII + noise

In addition to obscuring PII columns, we can deliberately add noise to other attributes to reduce the effectiveness of potential attacks. Differential privacy techniques, for instance, can provide formal guarantees against MIA.

Another approach involves randomly “swapping” data between entries. So for instance, in a demographic dataset, the ages of the included individuals might be reordered randomly in the records. This technique aims to provide plausible but randomized data by making it more difficult for an adversary to infer any information regarding any particular individual. These techniques aim to elevate privacy while preserving the utility of the data to a downstream task.

Depending on the amount of noise and the downstream task, some degree of utility degradation is expected.

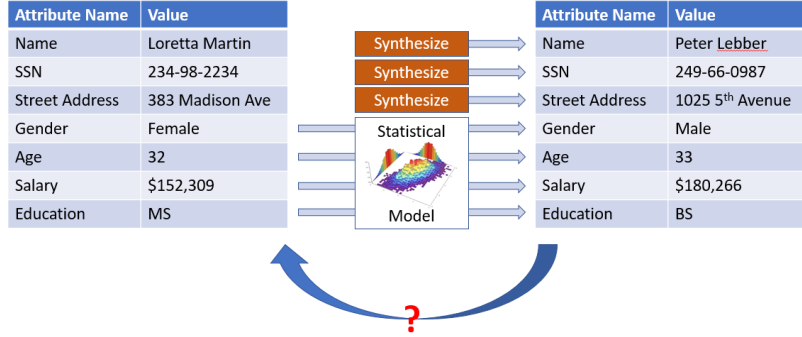


Figure 3: Privacy Level 3: Generative modeling. The question mark suggests the possibility of reverse-engineering the data.

### 4.3 Privacy Level 3: Generative modeling

Note that Privacy Levels 1 and 2 involve row-by-row transcription of the original data (with obfuscation or noise as appropriate). Accordingly, such datasets cannot be larger than the original.

With Level 3 we move to *generative* techniques where we analyze the original data to build a model that can create new data. Example approaches include Gaussian copula, and Generative-Adversarial-Networks (GAN) [PWV16, GPAM<sup>+</sup>14, PMG<sup>+</sup>18]. Other methods use differential privacy techniques to offer additional guarantees [ADR<sup>+</sup>19, XLW<sup>+</sup>18, YJvdS19]. In our own work, we have introduced a KD-tree-based formulation to model the data that offers additional protections as well [KNP<sup>+</sup>23].

All these methods enable the creation of new data elements distinct from the original data. They offer stronger protection than in Level 1 or Level 2, but are still potentially subject to attack. The risk is increased when the relative size of the generated data to the original data is large: For example, if we generate one million samples using an original dataset of only 1,000 we would expect to see generated samples clustering around the samples in the original data.

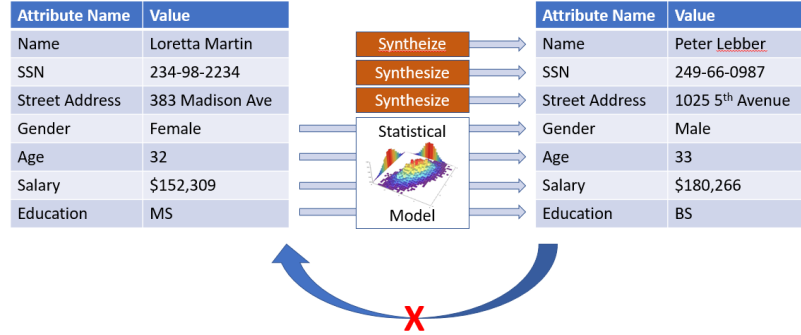


Figure 4: Privacy Level 4: Generative modeling + testing

#### 4.4 Level 4: Generative modeling + testing

For Level 4 we add explicit testing of each generated dataset to validate its resistance to specific attacks. The particular tests and the corresponding scores required to “pass” depend on the data and the application. For instance, it may be acceptable for certain properties of the data to “leak” while others should not. To operationalize this, we leverage published attack algorithms, then score the data depending on the success of the attack.

While it is hard to specify which test and which score would be necessary to achieve Level 4 privacy in all cases, the important and critical difference above Level 3, is the fact that the data is explicitly tested. The test and the scoring criteria must be determined by the individual business for the use case. Example scoring criteria measure resistance to membership inference, attribute reconstruction, and property attacks. among others [GBWT23, HCS<sup>+</sup>22, HJC<sup>+</sup>22, BDI<sup>+</sup>23, DL24].






Figure 5: Privacy Level 5: Calibrated simulation

#### 4.5 Level 5: Calibrated simulation

In this approach the generation method is not trained on real data. In fact, there is (usually) no learning in this approach. Instead, we rely on simulations governed by rules or knowledge of the process that would otherwise generate real data. These rules, however, are calibrated with reference to the real process such that the generated data follows some statistical properties of the original, real system. As an example, we might use a simulation of the stock market to generate stock price data. In our own work, we have developed calibrated simulations of equity markets that correspond to Level 5 privacy [VBP<sup>+</sup>19].

Utility degradation depends on the downstream task and the simulation framework. This approach generally represents a strong defense against adversarial attacks. However, they may be exposed to Property Inference Attacks, because the simulator is calibrated with respect to statistical properties of the real system.




Figure 6: Privacy Level 6: Uncalibrated simulation

#### 4.6 Level 6: Uncalibrated simulation

In this case we may not be aware of the statistical properties of the modeled system, or we might deliberately avoid adjusting the simulation to correspond to the properties of the original system. Even though such a simulation may not provide high fidelity data, it can still prove quite useful.

An important use is testing, in which we might use a simulation to explore all the potential values of data fields to see if they “break” our downstream processes. Additionally, we might choose to embed known examples of situations we want to be sure our systems detect (e.g., fraudulent transactions). Another use is to create what-if scenarios where we hypothesize the impact of one factor on another, to see if visualization techniques might enable us to discover those relationships in practice.

In general, this method yields a strong privacy guarantee. It remediates one of the consequences of level 5 generation of defence against PIA attacks, given that the statistical properties of the data is uncalibrated to the real dataset.

## 5 Summary

We describe six categories, or levels, of privacy protection for Financial Synthetic Data provided by different generation techniques. The strength of privacy protection relates to the resistance the technique offers against privacy attacks. Such attacks might enable an adversary to infer information about individual data points in the original data used to train a generator.

The six levels progress from least secure (Level 1) to most secure (Level 6), Level 1 depends on simple masking and obfuscation (which offers very little protection), while Level 6, uncalibrated simulation, provides the strongest protection. We focus specifically on financial data, but the categorizations may be useful in other industries (e.g. healthcare) and generation techniques as well.

## 6 Disclaimer

This paper was prepared for informational purposes by the CDAO group of JPMorgan Chase & Co and its affiliates (“J.P. Morgan”) and is not a product of the Research Department of J.P. Morgan. J.P. Morgan makes no representation and warranty whatsoever and disclaims all liability, for the completeness, accuracy or reliability of the information contained herein. This document is not intended as investment research or investment advice, or a recommendation, offer or solicitation for the purchase or sale of any security, financial instrument, financial product or service, or to be used in any way for evaluating the merits of participating in any transaction, and shall not constitute a solicitation under any jurisdiction or to any person, if such solicitation under such jurisdiction or to such person would be unlawful.

© 2024 JPMorgan Chase & Co. All rights reserved.

## References

- [ADM<sup>+</sup>20] Samuel A Assefa, Danial Dervovic, Mahmoud Mahfouz, Robert E Tillman, Prashant Reddy, Tucker Balch, and Manuela Veloso. Generating synthetic data in finance: opportunities, challenges and pitfalls. In *Proceedings of the First ACM International Conference on AI in Finance*, pages 1–8, 2020.
- [ADR<sup>+</sup>19] Hassan Jameel Asghar, Ming Ding, Thierry Rakotoarivelo, Sirine Mrabet, and Mohamed Ali Kaafar. Differentially private release of high-dimensional datasets using the gaussian copula, 2019.
- [BDI<sup>+</sup>23] Brian Belgodere, Pierre Dognin, Adam Ivankay, Igor Melnyk, Youssef Mroueh, Aleksandra Mojsilovic, Jiri Navartil, Apoorva Nitsure, Inkit Padhi, Mattia Rigotti, et al. Auditing and generating synthetic data with controllable trust trade-offs. *arXiv preprint arXiv:2304.10819*, 2023.
- [DL24] Yuntao Du and Ninghui Li. Towards principled assessment of tabular data synthesis algorithms. *arXiv preprint arXiv:2402.06806*, 2024.
- [GBWT23] Matteo Gioni, Franziska Boenisch, Christoph Wehmeyer, and Borbála Tasnádi. A unified framework for quantifying privacy risk in synthetic data, 2023.
- [GPAM<sup>+</sup>14] Ian J. Goodfellow, Jean Pouget-Abadie, Mehdi Mirza, Bing Xu, David Warde-Farley, Sherjil Ozair, Aaron Courville, and Yoshua Bengio. Generative adversarial networks, 2014.
- [HCS<sup>+</sup>22] Florimond Houssiau, Samuel N Cohen, Lukasz Szpruch, Owen Daniel, Michaela G Lawrence, Robin Mitra, Henry Wilde, and Callum Mole. A framework for auditable synthetic data generation. *arXiv preprint arXiv:2211.11540*, 2022.
- [HJC<sup>+</sup>22] Florimond Houssiau, James Jordon, Samuel N Cohen, Owen Daniel, Andrew Elliott, James Geddes, Callum Mole, Camila Rangel-Smith, and Lukasz Szpruch. Tapas: a toolbox for adversarial privacy auditing of synthetic data. In *NeurIPS 2022 Workshop on Synthetic Data for Empowering ML Research*, 2022.
- [KNP<sup>+</sup>23] Eleonora Kreačić, Navid Nouri, Vamsi K. Potluru, Tucker Balch, and Manuela Veloso. Differentially private synthetic data using KD-trees. In *The 39th Conference on Uncertainty in Artificial Intelligence*, 2023.
- [LWSF23] Zinan Lin, Shuaiqi Wang, Vyas Sekar, and Giulia Fanti. Summary statistic privacy in data sharing, 2023.

- [NS07] Arvind Narayanan and Vitaly Shmatikov. How to break anonymity of the netflix prize dataset, 2007.
- [PMG<sup>+</sup>18] Noseong Park, Mahmoud Mohammadi, Kshitij Gorde, Sushil Jajodia, Hongkyu Park, and Youngmin Kim. Data synthesis based on generative adversarial networks. *Proceedings of the VLDB Endowment*, 11(10):1071–1083, jun 2018.
- [PWV16] Neha Patki, Roy Wedge, and Kalyan Veeramachaneni. The synthetic data vault. In *2016 IEEE International Conference on Data Science and Advanced Analytics (DSAA)*, pages 399–410, 2016.
- [SSS16] Reza Shokri, Marco Stronati, and Vitaly Shmatikov. Membership inference attacks against machine learning models. *CoRR*, abs/1610.05820, 2016.
- [SZZ<sup>+</sup>23] Hui Sun, Tianqing Zhu, Zhiqiu Zhang, Dawei Jin, Ping Xiong, and Wanlei Zhou. Adversarial attacks against deep generative models on data: A survey. *IEEE Transactions on Knowledge and Data Engineering*, 35(4):3367–3388, apr 2023.
- [VBP<sup>+</sup>19] Svitlana Vyetrenko, David Byrd, Nick Petosa, Mahmoud Mahfouz, Danial Dervovic, Manuela Veloso, and Tucker Hybinette Balch. Get real: Realism metrics for robust limit order book market simulations, 2019.
- [XLW<sup>+</sup>18] Liyang Xie, Kaixiang Lin, Shu Wang, Fei Wang, and Jiayu Zhou. Differentially private generative adversarial network, 2018.
- [YJvdS19] Jinsung Yoon, James Jordon, and Mihaela van der Schaar. PATE-GAN: Generating synthetic data with differential privacy guarantees. In *International Conference on Learning Representations*, 2019.